

MASTERING AI AGENTS

A comprehensive guide to evaluating AI agents



Preface

In our previous e-book, “[Mastering RAG](#),” our goal was clear: building enterprise-grade RAG systems, productionizing them, monitoring their performance, and improving them. At the core of it, we understood how RAG systems enhance an LLM’s ability to work with specific knowledge by providing relevant context.

In this e-book, we’re taking a step further and asking, “How do we use LLMs to accomplish end-to-end tasks?” This singular question opens up a door: AI agents. A RAG system helps an LLM provide accurate answers based on given context. An AI agent takes that answer and actually does something with it — makes decisions, executes tasks, or coordinates multiple steps to achieve a goal.

A RAG-enhanced LLM could help answer questions about policy details by pulling relevant information. But an AI agent could actually process the claim end-to-end by analyzing the documentation, checking policy compliance, calculating payments, and even coordinating with other systems or agents when needed.

The ideas behind agents has existed for years. It can be a software program or another computational entity that can accept input from its environment and take actions based on rules. With AI agents, you’re getting what has never been there before: the ability to understand the context without predefined rules, the capacity to tune decisions based on context, and learning from every interaction. What you’re getting is not just a bot working with a fixed set of rules but a system capable of making advanced decisions in real-time.

Companies have quickly adapted, adopted, and integrated AI agents into their workflows. Capgemini’s research found that “10% of organizations already use AI agents, more than half plan to use them in 2025 and 82% plan to integrate them within the next three years.”

This e-book aims to be your go-to guide for all things AI agents. If you're a leader looking to guide your company to build successful agentic applications, this e-book can serve as a great guide to get you started. We also explore approaches to measuring how well your AI agents perform, as well as common pitfalls you may encounter when designing, measuring, and improving them.

The book is divided into five chapters:

Chapter 1 introduces AI agents, their optimal applications, and scenarios where they might be excessive. It covers various agent types and includes three real-world use cases to illustrate their potential.

Chapter 2 details three frameworks—LangGraph, Autogen, and CrewAI—with evaluation criteria to help choose the best fit. It ends with case studies of companies using these frameworks for specific AI tasks.

Chapter 3 explores the evaluation of an AI agent through a step-by-step example of a finance research agent.

Chapter 4 explores how to measure agent performance across systems, task completion, quality control, and tool interaction, supported by five detailed use cases.

Chapter 5 addresses why many AI agents fail and offers practical solutions for successful AI deployment.


We hope this book will be a great stepping stone in your journey to build trustworthy agentic systems.

- Pratik Bhavsar

Contents

Chapter 1: **What are AI agents**

7/27



Types of AI Agents	10
When to Use Agents?	21
When Not to Use Agents?	22
10 Questions to Ask Before You Consider an AI Agent	23
3 Interesting Real-World Use Cases of AI Agents	25

Chapter 2: **Frameworks for Building Agents**

28/43



LangGraph vs. AutoGen vs. CrewAI	30
Practical Considerations	31
What Tools and Functionalities Do They Support?	31
How Well Do They Maintain the Context?	32
Are They Well-Organized and Easy to Interpret?	33
What's the Quality of Documentation?	34
Do They Provide Multi-Agent Support?	34
What About Caching?	35
Looking at the Replay Functionality	35
What About Code Execution?	35
Human in the Loop Support?	37
Popular Use Cases Centered Around These Frameworks	40

Chapter 3:
How to Evaluate Agents**44/61**

Requirements	44
Defining the Problem	44
Define the React Agent	45
State Management	46
Create the Graph	47
Create the LLM Judge	54
Use Galileo Callbacks	55

Chapter 4:
**Metrics for Evaluating
AI Agents****62/79**

Case Study 1: Advancing the Claims Processing Agent	63
Case Study 2: Optimizing the Tax Audit Agent	66
Case Study 3: Elevating the Stock Analysis Agent	69
Case Study 4: Upgrading the Coding Agent	72
Case Study 5: Enhancing the Lead Scoring Agent	75

Chapter 5:
Why Most AI Agents Fail &
How to Fix Them

80/95

Development Issues	81
LLM Issues	82
Production Issues	86

01

CHAPTER

WHAT ARE AI
AGENTS?

What are AI agents?

Let's start by understanding what AI agents are and which tasks you should use them for to maximize their potential.

AI agents are software applications that use large language models (LLMs) to autonomously perform specific tasks, ranging from answering research questions to handling backend services. They're incredibly useful for tasks that demand complex decision-making, autonomy, and adaptability. You might find them especially helpful in dynamic environments where the workflow involves multiple steps or interactions that could benefit from automation.

[Salesforce estimates that salespersons spend 71% of their time on non-selling tasks \(like administrative tasks and manually entering data\)](#). Imagine the time that could have gone into directly engaging with customers, developing deeper relationships, and ultimately closing more sales. This is true across multiple domains and applications: finance, health care, tech, marketing, sales, and more.

Let's use an example to understand this better. Imagine you run an online retail business and receive hundreds of customer inquiries every day about order statuses, product details, and shipping information. Instead of answering each and every query yourself, you can integrate an AI agent into your solution to handle these queries.

Here's how it would typically work:

1. Customer Interaction

A customer messages your service asking, "When will my order ship?"

2. Data Retrieval

The AI agent accesses the order management system to find the specific order details.

3. Response Generation

Based on the data retrieved, the agent automatically provides an update to the customer, such as sending "Your order will ship tomorrow and you'll receive a tracking link via email once it's on its way."

The return to having an AI agent is multifold here:

- Super quick response time that keeps your customers happy
- Frees up your human staff to handle more complex queries and issues
- Improves your overall productivity and efficiency

Fig 1.1 is an example of how agents are leveraged for code generation.




Fig 1.1: Automated AI-Driven Development using AI agents

Types of AI Agents

Now that we're familiar with what AI agents are, let's look at different types of AI agents along with their characteristics, examples, and when you can use them.

See **Table 1.1** below to get a quick idea of the types of AI agents and where and when you can use them.

Name of the agent	Key Characteristics	Examples	Best For
Fixed Automation: The Digital Assembly Line	No intelligence, predictable behavior, limited scope	RPA, email autoresponders, basic scripts	Repetitive tasks, structured data, no need for adaptability
LLM-Enhanced: Smarter, but Not Einstein	Context-aware, rule-constrained, stateless	Email filters, content moderation, support ticket routing	Flexible tasks, high-volume/low-stakes, cost-sensitive scenarios
ReAct: Reasoning Meets Action	Multi-step workflows, dynamic planning, basic problem-solving	Travel planners, AI dungeon masters, project planning tools	Strategic planning, multi-stage queries, dynamic adjustments
ReAct + RAG: Grounded Intelligence	External knowledge access, low hallucinations, real-time data	Legal research tools, medical assistants, technical support	High-stakes decisions, domain-specific tasks, real-time knowledge needs
Tool-Enhanced: The Multi-Taskers	Multi-tool integration, dynamic execution, high automation	Code generation tools, data analysis bots	Complex workflows requiring multiple tools and APIs
Self-Reflecting: The Philosophers	Meta-cognition, explainability, self-improvement	Self-evaluating systems, QA agents	Tasks requiring accountability and improvement
Memory-Enhanced: The Personalized Powerhouses	Long-term memory, personalization, adaptive learning	Project management AI, personalized assistants	Individualized experiences, long-term interactions
Environment Controllers: The World Shapers	Active environment control, autonomous operation, feedback-driven	AutoGPT, adaptive robotics, smart cities	System control, IoT integration, autonomous operations
Self-Learning: The Evolutionaries	Autonomous learning, adaptive/scalable, evolutionary behavior	Neural networks, swarm AI, financial prediction models	Cutting-edge research, autonomous learning systems

Table 1.1: Types of agents and their characteristics

Fixed Automation – The Digital Assembly Line

This level of AI agents represents the simplest and most rigid form of automation. These agents don't adapt or think—they just execute pre-programmed instructions. They are like assembly-line workers in a digital factory: efficient but inflexible. Great for repetitive tasks, but throw them a curveball, and they'll freeze faster than Internet Explorer. (See **Table 1.2** below)

Feature	Description
Intelligence	No learning, adaptation, or memory.
Behavior	Predictable and consistent, follows pre-defined rules.
Scope	Limited to repetitive, well-defined tasks. Struggles with unexpected scenarios.
Best Use Cases	Routine tasks, structured data, situations with minimal need for adaptability.
Examples	RPA for invoice processing, email autoresponders, basic scripting tools (Bash, PowerShell).

Table 1.2: Characteristics of a fixed automation agent

The fixed automation workflow (See **Fig 1.2**) follows a simple, linear path. It begins when a specific input (like a file or data) triggers the system, which consults its predefined rulebook to determine what to do. Based on these rules, it executes the required action and finally sends out the result or output. Think of it as a digital assembly line where each step must be completed in exact order, without deviation.




Fig 1.2: Workflow of a fixed automation agent

LLM-Enhanced – Smarter, but Not Exactly Einstein

These agents leverage LLMs to provide contextual understanding and handle ambiguous tasks while operating within strict boundaries. [LLM-Enhanced Agents](#) balance intelligence and simplicity, making them highly efficient for low-complexity, high-volume tasks. Take a look at their features below in **Table 1.3**.

Feature	Description
Intelligence	Context-aware; leverages LLMs to process ambiguous inputs with contextual reasoning.
Behavior	Rule-constrained; decisions are validated against predefined rules or thresholds.
Scope	Stateless; no long-term memory; each task is processed independently.
Best Use Cases	Tasks requiring flexibility with ambiguous inputs, high-volume/low-stakes scenarios, and cost-sensitive situations where "close enough" is sufficient.
Examples	Email filters, AI-enhanced content moderation, customer support classification.

Table 1.3: Characteristics of an LLM-enhanced agent

The workflow below (**Fig 1.3**) shows how these smarter agents process information: starting with the input, the agent uses LLM capabilities to analyze and understand the input context. This analysis then passes through rule-based constraints that keep the agent within defined boundaries, producing an appropriate output. It's like having a smart assistant who understands context but still follows company policy before making decisions.




Fig 1.3: Workflow of a LLM-enhanced agent

ReAct – Reasoning Meets Action

ReAct agents combine Reasoning and Action to perform tasks that involve strategic thinking and multi-step decision-making. They break complex tasks into manageable steps, reasoning through problems dynamically and acting based on their analysis. These agents are like your type-A friend who plans their weekend down to the minute.

Table 1.4 lists their characteristics.

Feature	Description
Intelligence	Reasoning and action; mimics human problem-solving by thinking through a problem and executing the next step.
Behavior	Handles multi-step workflows, breaking them down into smaller, actionable parts. Dynamically adjusts strategy based on new data.
Scope	Assists with basic open-ended problem-solving, even without a direct solution path.
Best Use Cases	Strategic planning, multi-stage queries, tasks requiring dynamic adjustments, and re-strategizing.
Examples	Language agents solving multi-step queries, AI Dungeon Masters, project planning tools.

Table 1.4: Characteristics of a fixed ReAct agent

The ReAct workflow starts with an Input Query and then enters a dynamic cycle between the Reasoning and Action Phase, as you'll see in **Fig 1.4**. Unlike simpler agents, it can loop between thinking and acting repeatedly until the desired outcome is achieved before producing the final Output/Action. Think of it as a problem solver that keeps adjusting its approach - analyzing, trying something, checking if it worked, and trying again if needed.




Fig 1.4: Workflow of a ReAct agent

ReAct + RAG – Grounded Intelligence

Now, moving on to agents who are much more intelligent, we come to ReAct + RAG agents that combine reasoning, action, and real-time access to external knowledge sources. This integration allows them to make informed decisions grounded in accurate, domain-specific data, making them ideal for high-stakes or precision-critical tasks (especially when you add evaluations). These agents are your ultimate trivia masters with Google on speed dial. See **Table 1.5** to learn how this agent works.

Feature	Description
Intelligence	Employs a RAG workflow, combining LLMs with external knowledge sources (databases, APIs, documentation) for enhanced context and accuracy.
Behavior	Uses ReAct-style reasoning to break down tasks, dynamically retrieving information as needed. Grounded in real-time or domain-specific knowledge.
Scope	Designed for scenarios requiring high accuracy and relevance, minimizing hallucinations.
Best Use Cases	High-stakes decision-making, domain-specific applications, tasks with dynamic knowledge needs (e.g., real-time updates).
Examples	Legal research tools, medical assistants referencing clinical studies, technical troubleshooting agents.

Table 1.5: Characteristics of a ReAct + RAG agent

Starting with an Input Query, this advanced workflow combines ReAct's reasoning-action loop with an additional Knowledge Retrieval step. The agent cycles between Reasoning, Action Phase, and Knowledge Retrieval (See **Fig 1.5**) — consulting external sources as needed — until it reaches the desired outcome and produces an Output/Action. It's like having a problem solver who not only thinks and acts but also fact-checks against reliable sources along the way.




Fig 1.5: Workflow of a ReAct + RAG agent

Tool-Enhanced – The Multi-Taskers

Tool-enhanced agents are versatile problem solvers that integrate multiple tools, leveraging APIs, databases, and software to handle complex, multi-domain workflows. They combine reasoning, retrieval, and execution for seamless, dynamic task completion. Think of them as tech-savvy Swiss Army knives capable of combining reasoning, retrieval, and execution seamlessly! (See **Table 1.6**)

Feature	Description
Intelligence	Leverages APIs, databases, and software tools to perform tasks, acting as a multi-tool integrator.
Behavior	Handles multi-step workflows, dynamically switching between tools based on task requirements.
Scope	Automates repetitive or multi-stage processes by integrating and utilizing diverse tools.
Best Use Cases	Jobs requiring diverse tools and APIs in tandem for complex or multi-stage automation.
Examples	Code generation tools (GitHub CoPilot, Sourcegraph's Cody, Warp Terminal), data analysis bots combining multiple APIs.

Table 1.6: Characteristics of tool-enhanced agents

Starting with an Input Query, the agent combines reasoning with a specialized tool loop. After the initial reasoning phase, it selects the appropriate tool for the task (Tool Selection) and then executes it (Tool Execution). This cycle repeats until the desired outcome is achieved, leading to the final Output/Action. (See **Fig 1.6**)




Fig 1.6: Workflow of tool-enhanced agents

Self-Reflecting – The Philosophers

These agents think about their thinking. Self-reflecting agents introduce meta-cognition—they analyze their reasoning, assess their decisions, and learn from mistakes. This enables them to solve tasks, explain their reasoning, and improve over time, ensuring greater reliability and accountability. (See **Table 1.7**)

Feature	Description
Intelligence	Exhibits meta-cognition, evaluating its own thought processes and decision outcomes.
Behavior	Provides explanations for actions, offering transparency into its reasoning. Learns from mistakes and improves performance over time.
Scope	Suited for tasks requiring accountability and continuous improvement.
Best Use Cases	Quality assurance, sensitive decision-making where explainability and self-improvement are crucial.
Examples	AI that explains its reasoning, self-evaluating learning systems, quality assurance (QA) agents.

Table 1.7: Characteristics of self-reflecting agents

Starting with an Input Query, the agent goes through a cycle of Reasoning and Execution, but with a crucial additional step: Reflection. After each execution, it reflects on its performance and feeds those insights back into its reasoning process. This continuous loop of thinking, doing, and learning continues until the desired outcome is achieved, producing the final Output/Action. This is evident in **Fig 1.7**.




Fig 1.7: Workflow of self-reflecting agents

Memory-Enhanced – The Personalized Powerhouses

Give an agent a little memory, and you have the ultimate personal assistant. Memory-enhanced agents bring personalization to the forefront by maintaining historical context and remembering user preferences, previous interactions, and task history. They act as adaptive personal assistants, providing tailored experiences and continuous, context-aware support. These agents remember your preferences, track your history, and theoretically — would never (ever) forget your coffee order! (See Table 1.8)

Feature	Description
Intelligence	Possesses long-term memory, storing and recalling past interactions, preferences, and task progress.
Behavior	Provides context-aware personalization, adapting decisions and actions based on user-specific data and history. Learns and improves over time.
Scope	Excels at tasks requiring individualized experiences, tailored recommendations, and maintaining consistency across multiple interactions.
Best Use Cases	Personalized assistance, long-term interactions, tasks spanning multiple sessions.
Examples	Project management AI with task history, customer service bots tracking interactions, personalized shopping assistants.

Table 1.8: Characteristics of memory-enhanced agents

Look at **Fig 1.8:** Starting with an Input Query, this agent first recalls relevant past experiences and preferences (Memory Recall), then uses this context for Reasoning about the current task. After deciding on a course of action, it executes it (Action/Execution), updates its memory with new information (Memory Update), and produces the Output.




Fig 1.8: Workflow of memory-enhanced agents

Environment Controllers – The World Shapers

Environment-controlling agents extend beyond decision-making and interaction—they actively manipulate and control environments in real time. These agents are equipped to perform tasks that influence the digital landscape or the physical world, making them ideal for applications in automation, robotics, and adaptive systems. Think smart thermostats, but on steroids! (See **Table 1.9**)

Feature	Description
Intelligence	Autonomous learning; refines models and processes based on feedback, data, or environmental changes without manual updates.
Behavior	Adaptive and scalable, adjusting to changing conditions and new tasks. Exhibits evolutionary behavior, improving performance over time.
Scope	Suited for cutting-edge research and autonomous learning systems, offering high potential but requiring careful monitoring.
Best Use Cases	Situations where autonomous learning and adaptation are crucial, such as complex research, simulation, or dynamic environments.
Examples	Neural networks with evolutionary capabilities, swarm AI systems, autonomous robotics, financial prediction models.

Table 1.9: Characteristics of environment-controlling agents

Observe the workflow in **Fig 1.9** carefully. Starting with an Input Query, the agent first observes its surroundings (Perception Phase), reasons about the current state and required changes (Reasoning Phase), takes action to modify the environment (Action Phase), and then receives feedback about the changes (Feedback Phase). This cycle repeats until the desired goal is met, producing both an Output and changed system state.




Fig 1.9: Workflow of an environment-controlled agent

Self-Learning – The Evolutionaries

The holy grail of AI agents: those that can improve themselves over time. They learn, adapt, and evolve without needing constant babysitting. These agents improve themselves over time, learning from interactions, adapting to new environments, and evolving without constant human intervention. They combine elements of reasoning, memory, environment control, and self-reflection with autonomous learning capabilities to adapt and optimize their behavior.

Are they the future of AI? Potentially. Are they also terrifying? Without evaluations, observation, regulation, and oversight, very much so.

Feature	Description
Intelligence	Autonomous learning; refines models and processes based on feedback, data, or environmental changes without manual updates.
Behavior	Adaptive and scalable, adjusting to changing conditions and new tasks. Exhibits evolutionary behavior, improving performance over time.
Scope	Suited for cutting-edge research and autonomous learning systems, offering high potential but requiring careful monitoring.
Best Use Cases	Situations where autonomous learning and adaptation are crucial, such as complex research, simulation, or dynamic environments.
Examples	Neural networks with evolutionary capabilities, swarm AI systems, autonomous robotics, financial prediction models.

Table 1.10: Self-learning agents' characteristics

From the workflow in **Fig 1.10**, you'll realize how a self-learning agent are akin to an AI researcher that gets smarter with every experiment, constantly refining its methods and knowledge.

Starting with an Input Query, the agent enters a continuous cycle beginning with the Learning Phase where it processes available data, moves to Reasoning to analyze it, then takes Actions based on its analysis. The Feedback Phase evaluates results, leading to an Evolution Phase where the agent adapts and improves its models. This cycle repeats continuously, producing not just an Output but an evolved version of both the solution and the agent itself.




Fig 1.10: Workflow of a self-learning agent

What's fascinating is that each type has its own sweet spot—there's no “one-size-fits-all” solution. The key is matching the right agent type to your specific needs, whether you need the reliable consistency of fixed automation for routine tasks or the adaptive intelligence of self-learning agents for cutting-edge research.

When to Use Agents?

We've looked at the agent types and where each one excels. That said, you still need to be able to gauge where you'll need an AI agent. Agents are highly beneficial when tasks require complex decision-making, autonomy, and adaptability. They excel in environments where the workflow is dynamic and involves multiple steps or interactions that can benefit from automation. You'll see how workflows in different domains can benefit from the use of AI agents in **Table 1.11** below:

Domain	Task	Benefits of Using AI Agents
Customer Support	Handling queries, providing real-time assistance, issue escalation	Agents enhance the efficiency and customer experience by offering timely and accurate responses, allowing human staff to focus on more complex issues.
Research and Data Analysis	Gathering, processing, and analyzing data	They autonomously provide deep insights from large datasets, helping you understand patterns without manual effort.
Financial Trading	Real-time data processing	Agents excel in making quick decisions based on rapidly-changing market conditions.
Education	Personalized learning experiences	These agents adapt to each student's learning pace, offering tailored feedback and supporting unique learning journeys effectively.
Software Development	Code generation, debugging, and testing	Agents streamline the development process by handling repetitive tasks like coding and testing, improving code quality, and reducing development time. They also learn and improve over time, which continually enhances their assistance.

Table 1.11: Domains and applications that can benefit from the use of AI agents

When Not to Use Agents?

Agents offer many advantages, but there are certain scenarios in which deploying them might not be the best option.

If the tasks you're dealing with are straightforward, occur infrequently, or require only minimal automation, the complexity and cost of implementing AI agents might not make sense for you. Simple tasks that existing software solutions can handle efficiently do not necessarily benefit from the added intricacy of agent-based systems. In such cases, sticking with traditional methods can be more efficient and cost-effective.

Also, if your tasks require deep domain-specific knowledge or expertise—like conducting complex legal analyses, making intricate medical diagnoses, or handling high-stakes decision-making in unpredictable environments—these are typically better left to experienced professionals. When you rely solely on agents for these critical tasks, it can lead to suboptimal or even harmful outcomes.

That said, fields like psychotherapy, counseling, or creative writing thrive on the nuances of human emotion and the creative process—areas where agents largely fall short. In these domains, the human touch is irreplaceable and essential for achieving meaningful outcomes.

Implementing agents also requires a significant investment from you in terms of time, resources, and expertise. If you're running a small business or managing a project with a tight budget, the costs of developing and maintaining these agents may not justify the benefits. In highly regulated industries, your use of agents might be restricted due to compliance and security concerns as well, and ensuring agents adhere to stringent regulatory requirements can be very challenging and resource-intensive.

10 Questions to Ask Before You Consider an AI Agent

Before you consider using AI agents, you'll need to ask yourself a set of questions to help you evaluate if it's actually worth the time, capital, and resources you'll be putting into it:

01 **What is the complexity of the task?**

Is the task simple and repetitive, or does it involve complex decision-making that could benefit from automation?

02 **How often does the task occur?**

Is this a frequent task where automation could save significant time and resources, or is it a rare event that might not justify the investment?

03 **What is the expected volume of data or queries?**

Will the agent be handle large volumes of data or queries where speed and efficiency are crucial?

04 **Does the task require adaptability?**

Are the conditions under which the task is performed constantly changing, requiring adaptive responses that an AI can manage?

05 **Can the task benefit from learning and evolving over time?**

Is there a benefit to having a system that learns from its interactions and improves its responses or strategies over time?

06 **What level of accuracy is required?**

Is it critical that the task is performed with high accuracy, such as in medical or financial settings, where AI might need to meet high standards?

07 Is human expertise or emotional intelligence essential?

Does the task require deep domain knowledge, human intuition, or emotional empathy that AI currently cannot provide?

08 What are the privacy and security implications?

Does the task involve sensitive information that must be handled with strict privacy and security measures?

09 What are the regulatory and compliance requirements?

Are there specific industry regulations or compliance issues that need to be addressed when using AI?

10 What is the cost-benefit analysis?

Does the return on investment in terms of time saved, efficiency gained, and overall performance outweigh the costs of implementing and maintaining an AI system?

Take time to evaluate these questions; this will help you better determine if an AI agent fits your needs and how it could be effectively implemented to enhance your operations or services.

3 Interesting Real-World Use Cases of AI Agents

Now that we've learned what agents are and when to and when not to use them, it's time to go through some interesting real-world use cases of AI agents.

1. Wiley and Agentforce

Company:

Wiley

AI Agent:

Agentforce by Salesforce

Use Case:

Customer service automation

Problem:

Wiley faced challenges handling spikes in service calls during peak times, particularly at the start of new semesters when thousands of students use Wiley's educational resources.

Need:

The company needed an efficient customer service system to manage the increased volume and maintain positive customer experiences.

Solution:

Wiley invested in Salesforce's Agentforce, an AI agent designed to enhance customer service operations. This integration has significantly improved case resolution rates and faster resolution of customer queries, especially during peak times, such as the start of new semesters when demand spikes.

ROI:

A 40%+ increase in case resolution compared to their previous chatbot, a 213% ROI, and \$230K in savings

2. Oracle Health and Clinical AI agent

Company:

Oracle Health

AI Agent:

Clinical AI Agen

Use Case:

Enhancing patient-provider interactions

Problem:

Healthcare providers faced documentation and time management challenges during patient visits, leading to burnout and reduced patient engagement.

Need:

There was a need for a solution that could streamline clinical workflows and improve documentation accuracy while allowing providers more time to interact with patients.

Solution:

Oracle Health developed its Clinical AI Agent, which automates documentation processes and enhances patient-provider interactions through a multimodal voice user interface. This allows providers to access patient information quickly and generate accurate notes efficiently.

ROI:

AtlantiCare, using the Clinical AI Agent, reported a 41% reduction in total documentation time, saving approximately 66 minutes per day, which translates to improved productivity and enhanced patient satisfaction.

3. Magid and Galileo

Company:

Magid

AI Agent:

RAG-based system
powered with real-time
observability capabilities

Use Case:

Empowering newsrooms
with generative AI technology

Problem:

Magid, a leader in consumer intelligence for media brands, needed to ensure consistent, high-quality content in a fast-paced news environment. The complexity of diverse topics made it challenging to uphold accuracy, and errors could potentially lead to significant repercussions.

Need:


A robust observability system was essential for monitoring AI-driven workflows and ensuring the quality of outputs across various clients. This scalability was crucial for managing the daily production of numerous stories.

Solution:

[Magid integrated Galileo's real-time observability](#) capabilities into their product ecosystem. This integration provided production monitoring, relevant metrics for tracking tone and accuracy, and customization options tailored to Magid's needs.

ROI:

With Galileo, Magid achieved 100% visibility over inputs and outputs, enabling customized offerings as they scale. This visibility helps identify trends and develop client-specific metrics, enhancing the accuracy of news delivery.



We'll look at many more use cases across multiple domains throughout the rest of this e-book. We'll examine how agents have driven greater productivity, quicker resolutions, and helped things get done faster.

In the next chapter, we're going to learn features of three prominent frameworks for building AI agents. Lots of exciting stuff ahead!

Download the
full ebook
using the link below

